



PostgreSQL in a nutshell

ALL THE THINGS MATTER(?)

Disclaimer



- I am not a PostgreSQL specialist.
- This presentation should not be taken as an absolute source of truth.
- Everything here is a compilation of different knowledge sources.
- Last but not least: we're only scratching the surface!

What matters?

- PostgreSQL Architecture
 - Why the server side matters
 - Server architecture overview
 - Why statistics matter
 - Hands-on
 - Playground
 - Statistics & ANALYZE
 - VACUUM vs VACUUM FULL
 - EXPLAIN
- Indexing
 - Concept
 - Index types
 - Hands-on
 - B-tree
 - GIN
- JSON Support
 - History & motivation
 - Types
 - JSON
 - JSONB
 - JSONPATH
 - TOAST
 - Hands-on
 - JSONB

Why the server side matters



Do you remember the last time you installed a PostgreSQL database and configured it for your application?



Why does it matter?



PostgreSQL's default configuration is not suitable for production.



At the very least, buffer settings, max connections, and logging are things you should review.




Up-to-date statistics allow the system to make better decisions — such as choosing the correct execution plan.




Processes such as Autovacuum are crucial to keeping tables updated and providing accurate information to the query planner.


Server architecture overview



Server architecture overview



Server architecture overview



Server architecture overview

Clients(multiple types)

Query Processor

Postmaster
(process manager)

Connection Manager

Shared Buffer

Parser

Rewriter

Planner

Executor

PostgreSQL Server

Query Processor

Connection Manager

Parser

Rewriter

Planner

Executor

Utility Processes

Archiver

Check Pointer

Collector

Stats Collector

Vacuum



Data Files



WAL Files




Log Files





Archive Files



Storage Manager





Server architecture overview





Server architecture overview




Server architecture overview



Server architecture overview



Server architecture overview



Executor (Query Executor)


Postmaster

It is responsible for physically executing the execution plan generated by the Query Planner/Optimizer.

The **execution plan** is a tree of nodes (seq scan, index scan, hash join, sort, aggregate, limit, etc.).

Responsibilities:

- Execute each node of the plan
- Manage the flow of tuples
- Apply filter conditions
- Manage buffers and interact with the Storage Manager
- Enforce MVCC rules
- Trigger firing
- Produce results for the client



Server architecture overview

Clients(multiple types)

Utility Processes

Postmaster
(process manager)

Shared Buffer

Memory

Temp Tables

WAL Log

Index

Relational

Tablespaces

Filespace

File

Page

Block

Page

Server architecture overview

Utility Processes

Writer

WAL Writer

Archiver Process

Check Pointer

Logging Collector

Stats Collector

Vacuum

Checkpointer Process

Ensures periodic persistence of data.

- Writes "dirty" pages from shared buffers to data files.
- Reduces recovery time after a crash.
- Operates based on checkpoint_timeout, among other parameters.

Essential role:

Minimize recovery effort and ensure on-disk consistency.

Background Writer

Gradually writes modified pages to disk, avoiding write bursts during checkpoints.

- Smooths out I/O over time.

Essential role:

Keep the buffer pool healthy and reduce latency for backends.

WAL Writer

Responsible for writing data to the Write-Ahead Log.

- Ensures the core durability principle of ACID.
- Backends do not write to WAL directly — they send buffers to the writer.

Essential role:

Reduce contention and ensure WAL is written efficiently and in order.

Auto Vacuum Launcher



Orchestrates autovacuum processes.

- Monitors tables, detects bloat, and triggers vacuum operations as needed.

Essential role:

Preserve space, update statistics, and keep MVCC functioning.

Server architecture overview



Why statistics matters

PostgreSQL statistics are one of the essential pillars of database performance.

Without good statistics, the Query Planner literally goes blind, and the consequences are:

- Bad plans that can result in unnecessary seq scans.
- Incorrect index usage.
- Extremely expensive nested loops.
- Inefficient joins, massive sorts, memory explosions, and ultimately, timeouts

In short: a nightmare!

Why statistics matters



PostgreSQL uses a **cost-based query planner**.



Which means:

It chooses an execution plan based on cost estimates.
These estimates rely almost entirely on statistics.



So:

Good statistics → accurate estimates → good plans → fast queries
Bad statistics → wrong estimates → bad plans → slow queries



Why statistics matters



Statistics are produced in **two ways**:



Automatically via
Autovacuum

Autovacuum triggers
ANALYZE on tables that
have undergone enough
changes.



Manually, using the
ANALYZE command

For example, when we
run the command:
· ANALYZE my_table;
· ANALYZE my_table
(column1, column2)

Hands-on

Our playground

- Enough theory — let's set up our environment.
- What you'll need:
 - Docker
 - A simple IDE — I'll use VS Code. Sometimes I'll run code directly in it, and sometimes I'll use the terminal to send commands to PostgreSQL.
 - VS Code extensions:
 - <https://marketplace.visualstudio.com/items?itemName=ms-ossdata.vscode-pgsql>
 - <https://marketplace.visualstudio.com/items?itemName=inferrinizzard.prettier-sql-vscode>
- To make things easier, clone the presentation repository and follow the instructions:
<https://github.com/isaacvitor/postgresql-nutshell>

Hands-on

Statistics, ANALYZE, VACUUM and VACUUM FULL

- New statistics are generated when Autovacuum runs the Vacuum process together with ANALYZE.
- Autovacuum is an automated process that uses Vacuum. It runs automatically when certain thresholds are reached.
- Autovacuum settings can be changed in the *postgresql.conf* file.
- What is the Vacuum?
 - In summary, Vacuum is the process responsible for cleaning up dead tuples.
- There are two versions of the Vacuum process:
 - VACUUM – Marks spaces as being available for reuse. This process does not lock the table.
 - VACUUM FULL – Removes the deleted or updated records and reorders the table data. Note: VACUUM FULL requires an exclusive lock on the table.

Hands-on

Statistics, ANALYZE, VACUUM and VACUUM FULL

Let's see them in action!

- Statistics and ANALYZE
- VACUUM
- VACUUM FULL

Hands-on

EXPLAIN command

What is the EXPLAIN command?

- The EXPLAIN command displays the **execution plan** — the strategy PostgreSQL will use to execute a query. It does **not** execute the query itself (unless you use EXPLAIN ANALYZE)
- In summary, EXPLAIN shows us:
 - which algorithms will be used (Seq Scan, Index Scan, Bitmap Scan, Hash Join, Merge Join, etc.);
 - the order in which operations will occur;
 - cost estimates, number of rows, and row width;
 - use of parallelism;
 - filtering, join conditions, rechecks;
 - which indexes were chosen;
 - how the hierarchical execution plan is structured.

Hands-on

EXPLAIN command

Basic syntax

- EXPLAIN [ANALYZE] [VERBOSE] [COSTS] [SETTINGS] [BUFFERS] [WAL] [TIMING] [SUMMARY] [FORMAT format] query

Main parameters:

- EXPLAIN query — Basic plan
- EXPLAIN ANALYZE query — Executes the query and shows actual times
- EXPLAIN VERBOSE query — Detailed information
- EXPLAIN (ANALYZE, BUFFERS) — Includes buffer usage
- EXPLAIN (ANALYZE, VERBOSE, BUFFERS) — Complete information
- EXPLAIN (ANALYZE true, TIMING false) — Disables timing for faster queries

Output formats:

- EXPLAIN (FORMAT TEXT) — Default text format
- EXPLAIN (FORMAT JSON) — JSON
- EXPLAIN (FORMAT XML) — XML
- EXPLAIN (FORMAT YAML) — YAML

Useful combinations:

EXPLAIN (ANALYZE true, BUFFERS true, TIMING true, COSTS true)
query;

Let's see it in action!

Hands-on

EXPLAIN command

Running the command:

```
EXPLAIN SELECT * FROM explain_test WHERE id = 500000;
```

We'll get something like this:

```
Gather (cost=1000.00..11614.43 rows=1 width=14)
Workers Planned: 2
-> Parallel Seq Scan on explain_test (cost=0.00..10614.33 rows=1 width=14)
  Filter: (id = 500000)
(4 rows)
```

What the result means:

Gather – Coordinator node that collects results from parallel workers

cost=1000.00..11614.43 – (startup cost .. total cost)

rows=1 – Estimated to return 1 row

width=14 – 14 bytes per row on average

Workers Planned: 2 – PostgreSQL will use 2 parallel worker processes

Parallel Seq Scan – Parallel sequential scan on the table

cost=0.00..10614.33 – Cost of the parallel scan

Filter: (id = 500000) – Condition applied during the scan

(4 rows) – Quantity of rows in the output



Seq Scan? What?

Indexing

Concept



There's no substitute for a missing index. To achieve **good performance**, you need to have **proper indexing**.



To understand this statement, we must grasp the underlying concepts:

An index is an auxiliary **data structure** that **speeds up** the lookup of rows in tables, avoiding full table scans (sequential scan). It works as a "shortcut" to find specific values without reading all the table's blocks.

Indexing

Concept



Why do indexes matter in PostgreSQL?

PostgreSQL stores tables as heap files by default, with no physical ordering.



Without
indexes, any:

filter (WHERE)
lookup for a specific value
key-based JOIN
ORDER BY or GROUP BY



...will likely result in a Sequential Scan

Indexing

Concept



Indexes are not free!

An index is maintained in its own data structure, independent from the base table.



How do you choose
the right index?



Operators(=, LIKE, >,<, !=, @>, <@, etc)

Data type (TEXT, JSONB, Array)

Scenario (Combination of multiple data types)

Indexing

Index types



PostgreSQL provides 6 index types:



B-tree, Hash, GiST, SP-GiST, GIN, and BRIN



The most commonly used types are:

B-tree – default index type

GIN – great for arrays, JSONB, full-text-like structures

GiST – flexible, extensible, geometric, ranges

BRIN – ideal for very large, naturally ordered tables

Hands-on

B-tree (Balanced multi-way search tree)

How to create it:

```
CREATE INDEX idx_name  
ON table_name (column_name);
```

Best for:

- Equality lookups
- Range queries
- Sorting (ORDER BY), grouping
- Most typical queries
- Foreign keys

How it works:

- Balanced tree structure
- Keeps keys in sorted order
- Supports $O(\log n)$ lookups

Common operators supported:

- =
- <, <=, >, >=
- BETWEEN
- IS NULL, IS NOT NULL
- Pattern matching with left-anchored LIKE:
 - LIKE 'abc%'
 - ILIKE 'abc%' (with proper operator class)
- Supports ordering operators for ORDER BY

Hands-on

GIN (Generalized Inverted Index)

How to create it:

```
CREATE INDEX idx_name  
ON table_name  
USING gin (column_name);
```

Best for:

- Array fields
- JSONB data
- Full-text search (tsvector)
- Rows with multiple components (composite values)

Operators for arrays:

- && — overlap
- @> — array contains
- <@ — array is contained by
- = — equality

Operators for JSONB:


- ? — key exists
- ?| — any of these keys exist
- ?& — all keys exist
- @> — contains
- <@ — contained by

Operators for full-text (tsvector):

- @@ — match
- @@@ — phrase search (extensions)

JSON Support

History & motivation



JSON Support

JSON Type



What it is:

textual (string) storage
with JSON validity
checks on insert.



Properties:

Preserves whitespace
and exact key order as
inserted

No efficient native
indexing on content



Use cases:

When you must preserve
the exact JSON text
(formatting)

When you will never
query based on JSON
content



Limitations:

Expensive queries and
extraction (re-parsed on
every access)

No GIN support for
inner-structure queries

JSON Support

JSONB type



What it is:

binary, normalized representation of JSON (key order not preserved, no whitespace). Designed for search/querying. Introduced in 9.4.



Pros:

- Indexable with GIN/GIST (queries using @>, ?, ?|, ?&, etc.)
- Much faster search and extraction operations compared to json
- Supports efficient functions/operators (->, ->>, #>, #>>, jsonb_set, jsonb_insert, etc.)




Cons:

- Updates that change the whole document are costly (full rewrite in most cases)
- It is TOASTable — large documents may incur I/O overhead and degraded key-search performance when TOASTed

JSON Support

JSONPATH


- What it is:
 - a type for storing SQL/JSON Path expressions; introduced together with SQL/JSON Path support (Postgres 12).
Provides a standards-based JSONPath querying language. Useful with @?, @@, and jsonb_path_query* functions.



JSON Support

TOAST

- PostgreSQL stores data in **8 KB pages**.
When a value (such as JSONB, text, bytea, arrays) becomes **too large to fit comfortably inside the page**, PostgreSQL activates **TOAST** — *The Oversized-Attribute Storage Technique*.
- Large JSONB objects are one of the most common triggers for TOAST



JSON Support

TOAST

What TOAST does:

- When a column exceeds roughly 2 KB, PostgreSQL may:
 - Compress the value
 - Store it externally in a TOAST table
 - Keep only a lightweight pointer in the main row

Why PostgreSQL uses TOAST:

- Prevent oversized rows from splitting across multiple pages
- Keep hot pages small (better caching)
- Avoid slow sequential reads due to bloated tuples
- Reduce excessive WAL generation during updates

JSON Support

TOAST

- Important consequences
 - Even if a query needs **only one field** inside a large JSONB document, PostgreSQL may need to **fetch and decompress the entire TOASTed value**.
 - Updates like `jsonb_set` often **rewrite the whole document**, which is expensive.
 - Queries that filter using `jsonb -> 'key'` may still force a full fetch + decompress of 200 KB
 - A single `jsonb_set()` on a 500 KB document generates 500 KB of WAL



JSON Support

TOAST

- Popular mistake:
 - Put everything inside a JSONB column
- CREATE TABLE **bad_small** (pk serial primary key, jb jsonb);
- CREATE TABLE **good_small** (pk serial primary key, **id int**, jb jsonb);

Hands-on

JSONB – Operators performance


- Nested containers performance
- How different JSONB navigation operators behave in terms of performance
 - Sample table with nested objects of various sizes (**1 KB – 1 MB**) and various nesting levels (0–9), containing one short key and one long key.
 - Test expressions (operators being benchmarked)
 - **-> (arrow):**
jb -> 'obj' -> 'obj' -> ... -> 'obj' -> 'key'
 - **#> (path):**
jb #> '{obj,obj,...,obj,key}'
 - **[] (subscript):**
jb['obj'][['obj']]...[['obj']]['key']
 - **JSONPath query:**
jsonb_path_query_first(jb, '\$.obj.obj.... .obj.key')

Note: The graphs presented here are an attempt to reproduce the experiments shown at PGConf 2021.

Hands-on JSONB – Operators performance

Conclusions:

- 1. JSONB access cost increases with document size
 - Extraction operations become significantly slower once JSONs exceed ~50–100 KB, and the effect becomes very pronounced near 1 MB.
- 2. Nesting depth has a direct and measurable impact
 - Deeper JSON structures (levels 6–9) consistently incur higher latency than shallow ones, even with the same overall document size.
- 3. The **-> operator** is the most sensitive to size and depth
 - It shows sharp performance degradation, reaching over 1000 μ s for large and deeply nested documents.
- 4. The **#> operator** delivers the best overall performance
 - Its growth curve is smoother and remains stable even for large JSONs.
 - It is the most reliable operator for accessing nested structures efficiently.





Hands-on

JSONB – Operators performance

Conclusions:

- 5. The `[] subscript` behaves essentially the same as `->`
 - Since `[]` is syntactic sugar for `->`, it inherits the same scaling issues and becomes slow on large nested objects.
- 6. **JSONPath** provides flexibility, but at a higher cost
 - `.jsonb_path_query_first()` is consistently slower than direct operators.
 - It is great for expressive queries but inefficient for simple lookups in large structures.
- 7. Operator choice **matters much more for large documents**
 - For small JSONs the difference between operators is minimal,
 - But in medium and large JSONs the differences grow by orders of magnitude.
- 8. For critical paths, prefer direct and explicit access methods
 - Whenever possible:
 - **use `#>` for nested lookup.**
 - **avoid JSONPath when the path is fixed:**
 - **minimize excessive nesting in JSON design**





Keep that in mind



ALL OPERATORS HAVE A COMMON OVERHEAD: DETOAST
TIME + JSONB ITERATION TIME



That's all folks

